

The background of the slide features several glowing orange jellyfish swimming against a dark blue background. The jellyfish have translucent bell shapes with bright orange and yellow patterns, and long, flowing tentacles. One large jellyfish is on the left, another large one is in the center, and a smaller one is on the right.


MICROSERVICE DECOMPOSITION PATTERNS

Sam Newman

O'REILLY®

Building Microservices

DESIGNING FINE-GRAINED SYSTEMS



Sam Newman

**Sam
Newman
& Associates**



@samnewman

NEW BOOK!

Monolith To Microservices.

Monolith To Microservices is a forthcoming book on system decomposition from O'Reilly

How do you detangle a monolithic system and migrate it to a microservices architecture? How do you do it while maintaining business-as-usual? As a companion to *Building Microservices*, this new book details a multiple approaches for helping you transition from existing monolithic systems to microservice architectures. This book is ideal if you're looking to evolve your existing systems, rather than just rewriting everything from scratch.

Topics include:


- Should you migrate to microservices, and if you should, how do you prioritise where to start
- How do you incrementally decompose an application
- Discusses multiple migration patterns and where they apply
- Delves into details of database decomposition, including the impact of breaking referential and transactional integrity, new failure modes, and more
- The growing pains you'll experience as your microservice architecture grows

[Read The Early Access Version!](#)

O'REILLY®

Monolith to Microservices


Evolutionary Patterns to Transform Your Monolith




Sam Newman

<https://samnewman.io/books/monolith-to-microservices/>


@samnewman








SINGLE PROCESS MONOLITH



SINGLE PROCESS MONOLITH


All code packaged
into a single process




SINGLE PROCESS MONOLITH

All code packaged
into a single process

All data stored in a
single database




“MODULAR” MONOLITH



“MODULAR” MONOLITH


The code is broken
into modules



“MODULAR” MONOLITH

The code is broken
into modules

Each module
packaged together
into a single process





“MODULAR” MONOLITH

The code is broken
into modules


Each module
packaged together
into a single process

Highly underrated
option






3RD PARTY MONOLITH



3RD PARTY MONOLITH


Could be on-prem
software, or a SaaS
product



3RD PARTY MONOLITH

Could be on-prem
software, or a SaaS
product

You have limited to
no ability to change
the core system




3RD PARTY MONOLITH

Could be on-prem software, or a SaaS product

You have limited to no ability to change the core system


You ***might*** have access to underlying storage...



3RD PARTY MONOLITH

Could be on-prem software, or a SaaS product


You have limited to no ability to change the core system




You ***might*** have access to underlying storage...

...or perhaps APIs


DISTRIBUTED MONOLITH




DISTRIBUTED MONOLITH




DISTRIBUTED MONOLITH




DISTRIBUTED MONOLITH



DISTRIBUTED MONOLITH




DISTRIBUTED MONOLITH



High cost of
change


DISTRIBUTED MONOLITH



High cost of change

Larger-scoped deployments

DISTRIBUTED MONOLITH




High cost of change

Larger-scoped deployments

More to go wrong

DISTRIBUTED MONOLITH



High cost of change

Larger-scoped deployments

More to go wrong

Release co-ordination



SAFE


SAFE

Watch out, there's a release train coming!

Ops Meta-Metrics


The Currency You Use to Pay For Change


John Allspaw
VP Operations
Etsy.com



<http://www.flickr.com/photos/warby/3296379139>

<http://www.slideshare.net/jallspaw/ops-metametrics-the-currency-you-pay-for-change-4608108>






**If you stick with a release train for too long, you'll
end up with a distributed monolith**


The Distributed Monolith

For when life isn't already complicated enough

Continuous Delivery


Release on demand









DOMAIN DRIVEN DESIGN



DOMAIN DRIVEN DESIGN








@samnewman





Order
Management


























**The problems you are trying to solve with
microservices will drive *how* you decompose them**


@samnewman

You need a clear idea of what you are trying to achieve with microservices


LOUDNESS

10






Monolith




Monolith




Production

Monolith



Production



You won't appreciate the true horror, pain and suffering of microservices until you're running them in production

“If you do a big bang rewrite, the only thing you’re certain of is a big bang”

- Martin Fowler (paraphrased)



<https://www.flickr.com/photos/pahudson/3224209758/>

**So migration patterns we use must allow for
incremental change**

**Patterns that allow for our architecture to
evolve, whilst still delivering features**



IMPLEMENTING A STRANGLER FIG PATTERN

IMPLEMENTING A STRANGLER FIG PATTERN

1. Asset capture:

Identify the functionality to move to a new microservice

IMPLEMENTING A STRANGLER FIG PATTERN

1. Asset capture:

Identify the functionality to move to a new microservice

2. Redirect calls

Intercept calls to old functionality,
and redirect to the new service

“MOVING” FUNCTIONALITY?

“MOVING” FUNCTIONALITY?

It might be copy and paste

“MOVING” FUNCTIONALITY?


It might be copy and paste

More likely is a total or partial rewrite


HTTP PROXY




HTTP PROXY




HTTP PROXY




HTTP PROXY




HTTP PROXY





HTTP PROXY




HTTP PROXY











Branch By Abstraction


BRANCH BY ABSTRACTION




BRANCH BY ABSTRACTION




BRANCH BY ABSTRACTION



BRANCH BY ABSTRACTION




BRANCH BY ABSTRACTION



BRANCH BY ABSTRACTION


1. Create abstraction point



BRANCH BY ABSTRACTION

1. Create abstraction point


2. Start work on new service implementation



BRANCH BY ABSTRACTION

1. Create abstraction point


2. Start work on new service implementation




BRANCH BY ABSTRACTION

1. Create abstraction point

2. Start work on new service implementation




BRANCH BY ABSTRACTION




1. Create abstraction point
2. Start work on new service implementation
3. Switch over

BRANCH BY ABSTRACTION




1. Create abstraction point
2. Start work on new service implementation
3. Switch over

BRANCH BY ABSTRACTION




1. Create abstraction point
2. Start work on new service implementation
3. Switch over
4. Cleanup

BRANCH BY ABSTRACTION




1. Create abstraction point
2. Start work on new service implementation
3. Switch over
4. Cleanup


WORKING WITH LEGACY CODE




PARALLEL RUN




PARALLEL RUN




PARALLEL RUN



PARALLEL RUN



PARALLEL RUN



GITHUB SCIENTIST

Scientist!

A Ruby library for carefully refactoring critical paths. [build](#) passing [coverage](#) 99%

How do I science?

Let's pretend you're changing the way you handle permissions in a large web app. Tests can help guide your refactoring, but you really want to compare the current and refactored behaviors under load.

```
require "scientist"

class MyWidget
  def allows?(user)
    experiment = Scientist::Default.new "widget-permissions"
    experiment.use { model.check_user?(user).valid? } # old way
    experiment.try { user.can?(:read, model) } # new way

    experiment.run
  end
end
```

Wrap a `use` block around the code's original behavior, and wrap `try` around the new behavior. `experiment.run` will always return whatever the `use` block returns, but it does a bunch of stuff behind the scenes:

<https://github.com/github/scientist>

GITHUB SCIENTIST

Scientist!

A Ruby library for carefully refactoring critical paths. [build](#) passing [coverage](#) 99%

How do I science?

Let's pretend you're changing the way you handle permissions in a large web app. Tests can help guide your refactoring, but you really want to compare the current and refactored behaviors under load.

```
require "scientist"

class MyWidget
  def allows?(user)
    experiment = Scientist::Default.new "widget-permissions"
    experiment.use { model.check_user?(user).valid? } # old way
    experiment.try { user.can?(:read, model) } # new way

    experiment.run
  end
end
```


Wrap a `use` block around the code's original behavior, and wrap `try` around the new behavior. `experiment.run` will always return whatever the `use` block returns, but it does a bunch of stuff behind the scenes:

Alternatives


- [daylerees/scientist](#) (PHP)
- [scientistproject/scientist.net](#) (.NET)
- [joealcorn/laboratory](#) (Python)
- [rawls238/Scientist4J](#) (Java)
- [tomiaijo/scientist](#) (C++)
- [trello/scientist](#) (node.js)
- [ziyasal/scientist.js](#) (node.js, ES6)
- [yeller/laboratory](#) (Clojure)
- [lancew/Scientist](#) (Perl 5)
- [lancew/ScientistP6](#) (Perl 6)
- [MadcapJake/Test-Lab](#) (Perl 6)
- [cwbriones/scientist](#) (Elixir)
- [calavera/go-scientist](#) (Go)
- [jelmersnoeck/experiment](#) (Go)
- [sophtchev/scientist](#) (Kotlin / Java)
- [junkpiano/scientist](#) (Swift)

<https://github.com/github/scientist>


ACCESSING DATA




ACCESSING DATA



ACCESSING DATA




ACCESSING DATA



Our new service needs
some data....



ACCESSING DATA




Our new service needs
some data....




What are our options?


DATA REUSE?




DATA REUSE?




DATA REUSE?




DATA REUSE?




AVOID SHARING DATABASES!




AVOID SHARING DATABASES!




AVOID SHARING DATABASES!




AVOID SHARING DATABASES!




AVOID SHARING DATABASES!



AVOID SHARING DATABASES!




AVOID SHARING DATABASES!




Information hiding!


EXPOSE DATA IN THE MONOLITH




EXPOSE DATA IN THE MONOLITH




EXPOSE DATA IN THE MONOLITH




EXPOSE DATA IN THE MONOLITH (CONT)




EXPOSE DATA IN THE MONOLITH (CONT)




EXPOSE DATA IN THE MONOLITH (CONT)




EXPOSE DATA IN THE MONOLITH (CONT)




MOVE DATA TO THE NEW SERVICE



MOVE DATA TO THE NEW SERVICE




MOVE DATA TO THE NEW SERVICE




MOVE DATA TO THE NEW SERVICE

Monolith needs to be changed to source data from the new service



MOVE DATA TO THE NEW SERVICE

Monolith needs to be changed to source data from the new service



DB Refactoring Patterns


The Addison-Wesley Signature Series

REFACTORING DATABASES

EVOLUTIONARY DATABASE DESIGN

SCOTT W. AMBLER

PRAMOD J. SADALAGE




*Forewords by Martin Fowler, John Graham,
Sachin Rekhi, and Dr. Paul Dorsey*




MARTIN FOWLER
John Graham
Sachin Rekhi
Dr. Paul Dorsey
Book SIGNATURE


JOINS ACROSS TABLES




JOINS ACROSS TABLES




JOINS ACROSS TABLES




JOINS ACROSS TABLES




JOINS ACROSS TABLES



JOINS ACROSS TABLES





JOINS ACROSS TABLES





Catalog

Finance









ID	Album
123	Death Polka Vol 4
456	Best Of Music

Line Items

SKU	Date	Amount
123	14/7/19	\$4.55
456	14/7/19	\$7.67

Ledger

Database

.....

ID	Album
123	Death Polka Vol 4
456	Best Of Music

Line Items

SKU	Date	Amount
123	14/7/19	\$4.55
456	14/7/19	\$7.67

Ledger

Database




ID	Album
123	Death Polka Vol 4
456	Best Of Music

Line Items

SKU	Date	Amount
123	14/7/19	\$4.55
456	14/7/19	\$7.67

Ledger



Database





ID	Album
123	Death Polka Vol 4
456	Best Of Music

Line Items

Catalog DB






MOVING JOINS TO SERVICE CALLS


More network hops - latency can suffer

Lost enforcement of data consistency


STATIC REFERENCE DATA




SPLITTING TABLES




SPLITTING TABLES





SPLITTING TABLES





SPLITTING TABLES













**We can find out how
much something
costs...**




**We can find out how
much something
costs...**




**...and can take
payment...**

**We can find out how
much something
costs...**



**...but can't check stock
levels!**



**...and can take
payment...**

Should we keep selling CDs in this situation?

CAP Theory

CAP Theory

In a partition, you have to tradeoff
between consistency and availability

Take the money?

Take the money?

Favouring availability over consistency

Don't take the money?

Don't take the money?

Favouring consistency over availability

**This doesn't have to be black
and white...**

**This doesn't have to be black
and white...**

Last view of stock: 1 hour ago

Items: 100

**This doesn't have to be black
and white...**

Last view of stock: 1 hour ago [Sell item](#)
Items: 100

**This doesn't have to be black
and white...**

Last view of stock: 1 hour ago [Sell item](#)
Items: 100

Last view of stock: 1 hour ago
Items: 1

This doesn't have to be black
and white...

Last view of stock: 1 hour ago [Sell item](#)
Items: 100

Last view of stock: 1 hour ago [Don't sell item](#)
Items: 1

This doesn't have to be black and white...

Last view of stock: 1 hour ago [Sell item](#)
Items: 100

Last view of stock: 1 hour ago [Don't sell item](#)
Items: 1

Last view of stock: 2 days ago
Items: 100

This doesn't have to be black and white...

Last view of stock:	1 hour ago	Sell item
# Items:	100	
Last view of stock:	1 hour ago	Don't sell item
# Items:	1	
Last view of stock:	2 days ago	Don't sell item
# Items:	100	

Memories, Guesses, and Apologies

Rate this article ★★★★☆

 Pat Helland May 15, 2007

 11  0  5

Well, here I am blogging on the bus with my newly installed Windows Live Writer!!!

This blog is a text version of a five minute "Gong Show" presentation I did at CIDR (Conference on Innovative Database Research) on Jan 8,2007.

All computing can be considered as: "Memories, Guesses, and Apologies". This is a personal opinion about how computers suck. Furthermore, it offers additional opinions about how we can take advantage of their sucki-ness. Lets dig into this...

Newton and Einstein


It used to be that we thought of computing as one big-ass mainframe. The database folks only thought about the database. Transactions (and transactional serializability) offered a crisp and clear perspective of how time marches forward uniformly. When working on transaction T(i), any other transaction T(j) can be perceived as occurring before T(i) or after T(i). If T(i) and T(j) are concurrently processed, the transaction system ensures that either order is correct without modifying the semantics. This offers a crisp and clear perspective of now. Time marches forward like a clock exactly as Newton envisaged his universe.

Nowadays, we have lots and lots of computers. Big ones, small ones, connected, disconnected, occasionally connected, etc. These computers each have their own perspective of time. When you see data, it is unlocked and an artifact of the past. Time is subjective with many different notions of now. This is very much the way Einstein revamped our understanding of the universe.

Moving to SOA is like moving from Newton's Universe to Einstein's Universe.

<https://blogs.msdn.microsoft.com/pathelland/2007/05/15/memories-guesses-and-apologies/>

THANKS!



@samnewman



The screenshot of the Sam Newman website shows the homepage. At the top, the navigation menu includes 'Home', 'About', a yellow-highlighted 'Talks' button, 'Podcast', 'Writing', and 'Contact'. The main heading 'Sam Newman.' is followed by a sub-section titled 'Talks & Workshops.' Below this, there is a brief description of the talks and a link to 'Find Out More'. To the right, there is a section titled 'Book!' featuring the cover of 'Building Microservices' by Sam Newman. Another section titled 'Video!' shows a thumbnail for a video titled 'Feature Branches And Toggles In A Post-GitHub World.'

Sam Newman.

Talks & Workshops.

Here are a list of the talks I am currently presenting. On request, I can present different topics or even my older talks. If you want me to present these topics at your conference or company, then please contact me. You can also see where I'll be speaking next on my [events](#) page.

What Is This Cloud Native Thing Anyway? ■

45min Talk

A talk exploring what the hell Cloud Native means

→ [Find Out More](#)

Feature Branches And Toggles In A Post-GitHub World. ■

1. Validate the integration
2. Write the code branch, fix it
3. Integrate safely

Book! ■

Building Microservices

Video! ■

<https://samnewman.io/>

@samnewman